

Predicting Bike Rentals

ASHISH PREMRAJ

11 JANUARY, 2019

Contents

1	Introduction	2
1.1	Problem Statement	2
1.2	Data	2
2	Methodology	3
2.1	Pre Processing	3
2.1.1	Exploratory Data Analysis	3
2.1.2	Missing Value Analysis	8
2.1.3	Outlier Analysis	9
2.1.4	Feature Selection	10
2.2	Model Development	12
2.2.1	Data Splitting	12
2.2.2	Multicollinearity	13
2.2.3	Multiple Linear Regression	13
2.2.4	Decision Trees	18
2.2.5	Random Forest	19
3	Conclusion	20
3.1	Model Evaluation	20
3.1.1	Mean Absolute Percentage Error (MAPE)	20
3.1.2	Mean Squared Error (MSE)	21
3.2	Model Selection	21
	Appendix A - R Code	22

Chapter 1

Introduction

1.1 Problem Statement

To increase the sales in renting bikes, this project aims at building an effective model to predict bike rental counts on daily basis based on environmental and seasonal settings. Based on some of the features like seasons, work day, weekends, holiday, weather etc. we would like to predict the days we can expect high rental counts, moderate and low rental counts. This would benefit the business in increasing the sales and make effective business decisions.

1.2 Data

Our task is to build a regression model to predict the number of bike rentals on a daily basis based on multiple features like seasons and weather. Given below is the sample data set that we will be using to predict our bike rental counts.

```
#Load data
data_df = read.csv("day.csv")
head(data_df)
```

##	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit
## 1	1	2011-01-01	1	0	1	0	6	0	2
## 2	2	2011-01-02	1	0	1	0	0	0	2
## 3	3	2011-01-03	1	0	1	0	1	1	1
## 4	4	2011-01-04	1	0	1	0	2	1	1
## 5	5	2011-01-05	1	0	1	0	3	1	1
## 6	6	2011-01-06	1	0	1	0	4	1	1

##	temp	atemp	hum	windspeed	casual	registered	cnt
## 1	0.344167	0.363625	0.805833	0.1604460	331	654	985
## 2	0.363478	0.353739	0.696087	0.2485390	131	670	801
## 3	0.196364	0.189405	0.437273	0.2483090	120	1229	1349
## 4	0.200000	0.212122	0.590435	0.1602960	108	1454	1562
## 5	0.226957	0.229270	0.436957	0.1869000	82	1518	1600
## 6	0.204348	0.233209	0.518261	0.0895652	88	1518	1606

As you can see from the table below we have 15 variables, using which we will correctly predict our target variable.

```
## [1] "instant"    "dteday"    "season"    "yr"        "mnth"
## [6] "holiday"    "weekday"    "workingday" "weathersit" "temp"
## [11] "atemp"      "hum"        "windspeed" "casual"     "registered"
```

Chapter 2

Methodology

2.1 Data Preprocessing


Data Preprocessing is a crucial step in data science projects to effectively evaluate our model. Data Preprocessing is a data mining technique that involves transforming raw data into a meaningful or clean data set. To do that we follow some data preprocessing steps like Exploratory Data Analysis, Outlier Analysis, Missing Value Analysis, Feature Engineering and Feature Selection.

2.1.1 Exploratory Data Analysis



Exploratory Data Analysis is an approach of analyzing data set to summarize their main characteristics, often with visual methods. It focuses more narrowly on checking assumptions required for model fitting and hypothesis testing, handling missing values and making transformations of variables as needed.

Here below we can see the distribution of our dependent variable - count


```
#Analysis on dependent variable - Count  
ggplot(data = data_df, aes(cnt))+  
  geom_histogram()
```



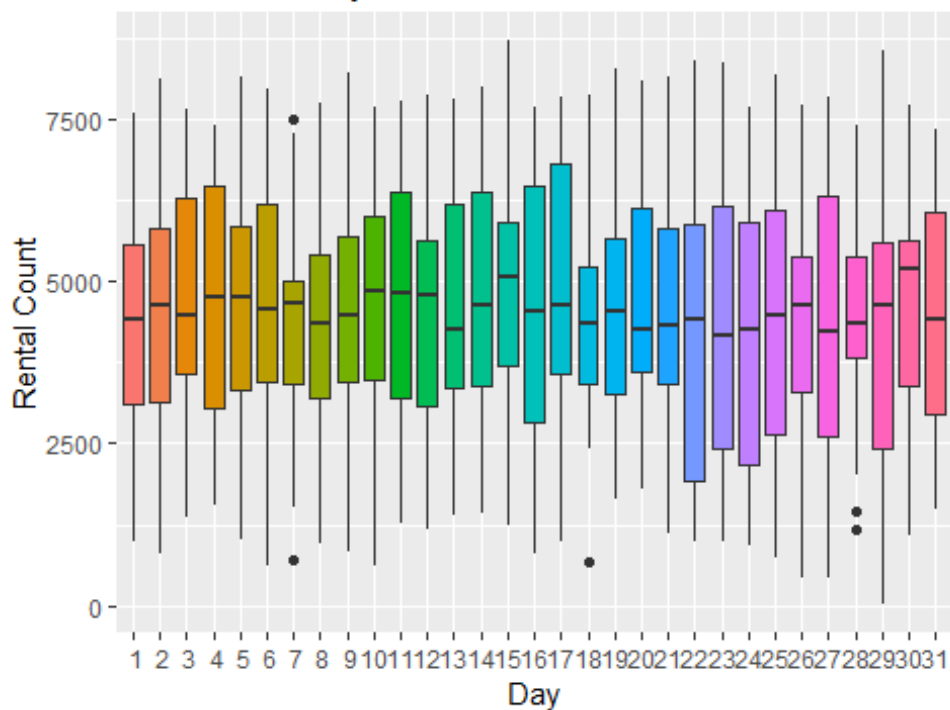
Analysis of all categorical independent variables vs dependent variable







Month wise Bike Rental counts



Day wise Bike Rental counts







From the above figures we can note the variation in rental counts for different seasons, month, day, year, weather conditions, weekdays and holidays. There were high rentals when the weather was clear with few clouds and low rentals during rain and thunderstorms. The rentals during 2012 was higher than 2011. We can also see the variations during different seasons will high rentals during fall.

Distribution of Continuous variables

```
gridExtra::grid.arrange(Hg1, Hg2, Hg3, Hg4, Hg5, Hg6)
```

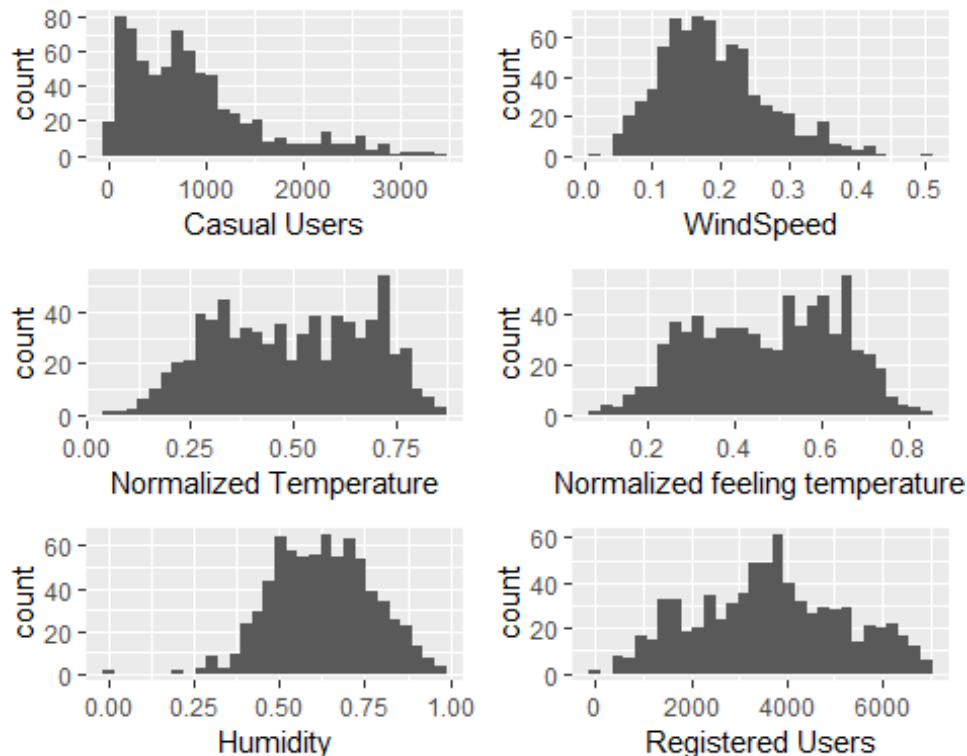


Figure 1.1 Distribution of continuous variables

From above figure, we see that most of the features are normally distributed. The Casual users is right-skewed with some outliers which we will analyze in below steps.

2.1.2 Missing Value Analysis

Dealing with missing values is one of the first data preprocessing steps. They need to be treated by imputation methods or by ignoring the entire feature if the percentage of missing values are high. There are different methods of imputation like mean, median, mode, KNN method. In our dataset, we do not have any missing values as we can see from below data frame.

```
Missing_val = data.frame(apply(data_df, 2, function(x){sum(is.na(x))}))
Missing_val$Columns = row.names(Missing_val)
row.names(Missing_val) = NULL
names(Missing_val)[1] = "Count"
```

```

#Display Missing Value counts - There are no missing values
Missing_val = Missing_val[c(2,1)]
Missing_val

##      Columns Count
## 1      instant     0
## 2       dteday     0
## 3        season     0
## 4         yr       0
## 5        mnth       0
## 6       holiday     0
## 7       weekday     0
## 8  workingday     0
## 9  weathersit       0
## 10        temp     0
## 11        atemp     0
## 12         hum      0
## 13  windspeed     0
## 14        casual     0
## 15 registered     0
## 16         cnt      0
## 17         day      0

```

2.1.3 Outlier Analysis

Outlier Analysis is a part of data preprocessing steps. The outliers in the data needs to be dealt with as they can lead to misleading results in our prediction. Different ways to deal with outliers are by removing the variables if they have a higher percentage of outliers or by imputation method. We can use boxplot.stats to analyze the possible outliers in the feature and then deal with them accordingly. From below fig 1.2, we can see the outliers on a few variables with casual users feature having high percentage of outliers.

```

numeric_data = data_df[,c("windspeed", "casual", "hum")]
cnames = colnames(numeric_data)
for(i in 1:length(cnames)){
  assign(paste0("Gn",i),
        ggplot(data_df, aes_string(y = cnames[i]))+
          stat_boxplot(geom = "errorbar", width = 0.5) +
          geom_boxplot(outlier.color = "red", outlier.shape = 18, outlier.
size = 2))
}

gridExtra::grid.arrange(Gn1,Gn2,Gn3,ncol=3)

```




Figure 1.2 Outlier Analysis using boxplot

2.1.4 Feature Selection

Correlation analysis will help us determine the relation between two or more variables. If we have features that are high correlated, we can consider only one of the features instead of taking all the correlated features as they produce redundant information and can cause misleading results.

From our analysis and figure 1.3, we can see that the features temp and atemp has high correlation which needs to be dealt with for better performance of our model. There are many ways in dealing with multicollinearity which we will discuss below while dealing with multicollinearity cases.

#Correlation Analysis

```
numeric_index = sapply(data_df, is.numeric)
numeric_data = data_df[,numeric_index]
train_cor = cor(numeric_data)
corrplot(train_cor, method = 'color', addCoef.col = 'black')
```




Figure 1.3 Correlation analysis

Before performing any type of modeling we need to assess the importance of each predictor variable in our analysis. There is a possibility that some of the features in our dataset do not add any information for our modeling. We need to get rid of those and some other features may have high correlation which can cause redundant information. So we need to deal with those as well by picking any one.

By using Random Forest algorithm we can categorize variables with high importance and consider them for our data modeling. From below figure 1.4, we see that feature yr has high importance compared to other features.

```
#Feature selection using Random Forest

RF_model = randomForest(cnt ~., data = data_df, ntree = 100, importance = T,
RUE)
pd = as.data.frame(importance(RF_model, type = 1))

#Plot features with imporatance from High to Low
ggplot(pd, aes(x = reorder(columns, -pd$`%IncMSE`), y = pd$`%IncMSE`, fill
= 'blue'))+
  geom_bar(stat = 'identity', show.legend = FALSE)+
  xlab("Variable Names")+
  ylab("Feature Importance Rate (%IncMSE)")+
  ggtitle("Features with Importance value")+
  theme(plot.title = element_text(hjust=0.5))
```


2.2.2 Multicollinearity

Multicollinearity occurs when two or more independent variables are highly correlated with each other which can lead to misleading results in building our model.

We can test the multicollinearity in our data using VIF(Variance Inflation Factor). Higher the value greater the variance. From below results using vifcor function we see that temp and atemp features have collinearity problem.

```
vif(numeric_data[, -5])

##   Variables      VIF
## 1      temp 62.969819
## 2      atemp 63.632351
## 3        hum  1.079267
## 4 windspeed  1.126768

vifcor(numeric_data[, -5], th = 0.9)

## 1 variables from the 4 input variables have collinearity problem:
##
## atemp
##
## After excluding the collinear variables, the linear correlation coefficients ranges between:
## min correlation ( hum ~ temp ):  0.1269629
## max correlation ( windspeed ~ hum ): -0.2484891
##
## ----- VIFs of the remained variables -----
##   Variables      VIF
## 1      temp 1.034283
## 2        hum 1.074850
## 3 windspeed 1.084580
```

2.2.3 Multiple Linear Regression

Linear Regression is a modelling approach to find relationship between the dependent variable and one or more independent variables. After applying linear regression algorithm on all features, we got a R-Squared value of 84%. We can further try to improve our model efficiency by performing some regularization techniques or stepwise model selection.

```
lm_model = lm(data = train, cnt ~ .)
summary(lm_model)

##
## Call:
## lm(formula = cnt ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3522.7  -370.8    54.3   458.9  2458.3
```

```
##
## Coefficients: (1 not defined because of singularities)
##
## Estimate Std. Error t value
## (Intercept) 1216.587 342.182 3.555
## seasonSummer 946.415 207.978 4.551
## seasonFall 784.411 236.093 3.322
## seasonWinter 1508.130 201.190 7.496
## yr2012 1987.106 65.006 30.568
## mnthFeb 69.421 160.543 0.432
## mnthMar 516.221 195.772 2.637
## mnthApr 386.944 288.709 1.340
## mnthMay 699.186 313.272 2.232
## mnthJun 414.511 325.024 1.275
## mnthJul -62.437 352.524 -0.177
## mnthAug 522.799 340.018 1.538
## mnthSep 995.448 294.087 3.385
## mnthOct 484.209 268.234 1.805
## mnthNov -99.993 254.062 -0.394
## mnthDec -90.872 200.966 -0.452
## holidayYes -671.729 197.743 -3.397
## weekdayMon 291.928 123.775 2.359
## weekdayTue 307.003 121.395 2.529
## weekdayWed 396.235 121.384 3.264
## weekdayThu 384.152 119.948 3.203
## weekdayFri 428.101 121.388 3.527
## weekdaySat 566.295 124.602 4.545
## workingdayYes NA NA NA
## weathersitMist + Cloudy -488.572 90.387 -5.405
## weathersitLight Snow + Rain + Thunderstorm -1767.697 233.122 -7.583
## temp 2550.835 1481.166 1.722
## atemp 2222.353 1531.693 1.451
## hum -1546.873 355.115 -4.356
## windspeed -2677.376 475.500 -5.631
## day 2 275.686 255.149 1.080
## day 3 282.892 253.341 1.117
## day 4 387.068 262.637 1.474
## day 5 193.747 272.175 0.712
## day 6 318.992 271.206 1.176
## day 7 50.645 256.180 0.198
## day 8 -7.757 268.307 -0.029
## day 9 178.327 264.293 0.675
## day10 376.196 253.065 1.487
## day11 461.379 255.882 1.803
## day12 305.921 256.628 1.192
## day13 263.576 265.544 0.993
## day14 435.662 272.958 1.596
## day15 446.758 257.710 1.734
## day16 422.559 256.354 1.648
## day17 722.665 265.811 2.719
## day18 111.797 255.922 0.437
```

## day19	298.371	255.445	1.168
## day20	398.096	260.916	1.526
## day21	355.254	257.038	1.382
## day22	-242.033	263.296	-0.919
## day23	116.829	264.991	0.441
## day24	51.891	255.657	0.203
## day25	-65.460	269.085	-0.243
## day26	285.966	262.178	1.091
## day27	160.820	272.594	0.590
## day28	98.299	259.227	0.379
## day29	-175.132	257.214	-0.681
## day30	-48.935	263.220	-0.186
## day31	552.304	303.343	1.821
##	Pr(> t)		
## (Intercept)	0.000411	***	
## seasonSummer	6.65e-06	***	
## seasonFall	0.000954	***	
## seasonWinter	2.80e-13	***	
## yr2012	< 2e-16	***	
## mnthFeb	0.665617		
## mnthMar	0.008615	**	
## mnthApr	0.180739		
## mnthMay	0.026042	*	
## mnthJun	0.202755		
## mnthJul	0.859487		
## mnthAug	0.124754		
## mnthSep	0.000765	***	
## mnthOct	0.071617	.	
## mnthNov	0.694053		
## mnthDec	0.651328		
## holidayYes	0.000733	***	
## weekdayMon	0.018711	*	
## weekdayTue	0.011730	*	
## weekdayWed	0.001168	**	
## weekdayThu	0.001444	**	
## weekdayFri	0.000457	***	
## weekdaySat	6.82e-06	***	
## workingdayYes	NA		
## weathersitMist + Cloudy	9.81e-08	***	
## weathersitLight Snow + Rain + Thunderstorm	1.53e-13	***	
## temp	0.085622	.	
## atemp	0.147398		
## hum	1.59e-05	***	
## windspeed	2.92e-08	***	
## day 2	0.280417		
## day 3	0.264655		
## day 4	0.141137		
## day 5	0.476873		
## day 6	0.240046		
## day 7	0.843362		


```

## day 8          0.976946
## day 9          0.500141
## day10          0.137729
## day11          0.071944 .
## day12          0.233764
## day13          0.321366
## day14          0.111070
## day15          0.083579 .
## day16          0.099878 .
## day17          0.006769 **
## day18          0.662405
## day19          0.243317
## day20          0.127667
## day21          0.167522
## day22          0.358387
## day23          0.659480
## day24          0.839236
## day25          0.807891
## day26          0.275888
## day27          0.555469
## day28          0.704690
## day29          0.496245
## day30          0.852589
## day31          0.069215 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 762.9 on 528 degrees of freedom
## Multiple R-squared:  0.8608, Adjusted R-squared:  0.8455
## F-statistic: 56.3 on 58 and 528 DF, p-value: < 2.2e-16

#Since we have multicollinearity issue between temp and atemp features we will regularize it
#Feature modeling using Stepwise modeling selection - both forward and backward
lm_AIC = stepAIC(lm_model, direction = 'both')

## Start: AIC=7847.74
## cnt ~ season + yr + mnth + holiday + weekday + workingday + weathersit +
##      temp + atemp + hum + windspeed + day
##
##
## Step: AIC=7847.74
## cnt ~ season + yr + mnth + holiday + weekday + weathersit + temp +
##      atemp + hum + windspeed + day
##
##
##           Df Sum of Sq      RSS      AIC
## - day       30  25287408 332555269 7834.2
## <none>                307267862 7847.7
## - atemp      1   1225084 308492946 7848.1

```

```

## - temp      1  1725999 308993861 7849.0
## - holiday   1   6715362 313983224 7858.4
## - weekday   6  13907175 321175036 7861.7
## - hum       1  11042135 318309997 7866.5
## - windspeed 1  18450198 325718060 7880.0
## - weathersit 2   37595239 344863101 7911.5
## - season    3   39495966 346763828 7912.7
## - mnth      11   50463459 357731321 7915.0
## - yr        1  543772441 851040302 8443.7
##
## Step: AIC=7834.16
## cnt ~ season + yr + mnth + holiday + weekday + weathersit + temp +
##      atemp + hum + windspeed
##
##           Df Sum of Sq      RSS      AIC
## - atemp      1      776074 333331343 7833.5
## <none>                                332555269 7834.2
## - temp       1     2836961 335392230 7837.1
## - holiday    1     6575627 339130897 7843.7
## - weekday     6    13992851 346548120 7846.4
## + day        30    25287408 307267862 7847.7
## - hum        1    14628101 347183371 7857.4
## - windspeed  1    19729030 352284299 7866.0
## - weathersit  2    36002002 368557271 7890.5
## - mnth       11    49532534 382087803 7893.7
## - season     3     40576712 373131981 7895.7
## - yr         1    560334204 892889473 8411.9
##
## Step: AIC=7833.53
## cnt ~ season + yr + mnth + holiday + weekday + weathersit + temp +
##      hum + windspeed
##
##           Df Sum of Sq      RSS      AIC
## <none>                                333331343 7833.5
## + atemp      1      776074 332555269 7834.2
## - holiday    1     6915190 340246533 7843.6
## - weekday     6    13933002 347264345 7845.6
## + day        30    24838397 308492946 7848.1
## - hum        1    14340767 347672110 7856.3
## - windspeed  1    22207712 355539055 7869.4
## - weathersit  2    37020589 370351932 7891.3
## - mnth       11    49080954 382412297 7892.2
## - season     3     40864355 374195698 7895.4
## - temp       1     64311871 397643214 7935.1
## - yr         1    559735938 893067281 8410.0

```

Stepwise model regression is a method of fitting regression models in which the choice of selecting predictor variables is carried by an automatic procedure. The model is evaluated based on AIC score. Lesser the score better the prediction rate of those variables considered.

In our analysis, we have used stepwise model selection on our linear regression model with both forward and backward approach which evaluates the model by iterating through each of the features and considering the best fit with score of AIC.

2.2.4 Decision Tree

Now, we will use a different regression model to predict our rental counts. We will use decision tree algorithm to predict the values of our target variable.

```
dt_model = rpart(cnt ~., data = train, method = 'anova')
dt_model

## n= 587
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
##  1) root 587 2207581000 4505.470
##    2) temp< 0.41875 227  497753300 2973.511
##      4) yr=2011 119  119353800 2124.571
##        8) season=Spring,Summer 83  25066270 1635.024 *
##        9) season=Winter 36  28535080 3253.250 *
##      5) yr=2012 108  198138100 3908.917
##        10) season=Spring 60  57924380 3152.817
##          20) atemp< 0.294789 33  21642160 2591.667 *
##          21) atemp>=0.294789 27  13190340 3838.667 *
##        11) season=Summer,Winter 48  63035970 4854.042
##          22) day= 9,22,23,24,25,27,30 9  13613860 3133.444 *
##          23) day= 1, 2, 3, 5, 6, 7, 8,10,11,12,13,14,15,16,17,18,19,20,21,2
##              6,28,29,31 39  16629380 5251.103 *
##      3) temp>=0.41875 360  841155500 5471.456
##        6) yr=2011 175  123368700 4269.377
##          12) mnth=Feb,Mar,Apr,Nov,Dec 35  20214490 3403.800 *
##          13) mnth=May,Jun,Jul,Aug,Sep,Oct 140  70375630 4485.771
##            26) hum>=0.886187 10  8787794 2938.700 *
##            27) hum< 0.886187 130  35812440 4604.777 *
##        7) yr=2012 185  225708300 6608.557
##          14) hum>=0.8095835 13  28133180 4633.308 *
##          15) hum< 0.8095835 172  143020700 6757.849
##            30) mnth=Jan,Mar,Apr,May,Jun,Jul,Nov,Dec 109  89141900 6476.523 *
##            31) mnth=Aug,Sep,Oct 63  30326460 7244.587 *
```

2.2.5 Random Forest

Random Forest is another algorithm that can be used for regression analysis. We will use `randomForest` on our train data set to predict our target variable.

```
rf_model = randomForest(cnt ~.,data = train, ntree = 100)
rf_model

##
## Call:
## randomForest(formula = cnt ~ ., data = train, ntree = 100)
##              Type of random forest: regression
##              Number of trees: 100
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 646323.1
##              % Var explained: 82.81
```

Chapter 3

Conclusion

3.1 Model Evaluation

3.1.1 Mean Absolute Percentage Error (MAPE)

MAPE or Mean Absolute Percentage Error is a measure of prediction accuracy. It is commonly used as a loss function for regression problems and in model evaluation. We will use this metric on our data models to evaluate the measure of prediction accuracy. MAPE can be calculated by taking the mean of (Actual value – Predicted value).

From below analysis, after predicting our built models on the test dataset, we then calculated the MAPE result. We see that linear regression algorithm with stepwise fitting regression model gave low prediction error of 19.93% and accuracy of 80.07%.

```
predict_lm = predict(lm_model, test[,c(-12)])

## Warning in predict.lm(lm_model, test[, c(-12)]): prediction from a rank-
## deficient fit may be misleading

predict_lm_AIC = predict(lm_AIC, test[, -12])
predict_dt = predict(dt_model, test[, -12])
predict_rf = predict(rf_model, test[, -12])

#Finding error metrics
#Calculating MAPE - y= Actual value, yhat = Predicted values
mape = function(y,yhat){
  mean(abs((y-yhat)/y))*100
}

mape(test[,12], predict_lm)

## [1] 20.16607

mape(test[,12], predict_lm_AIC)

## [1] 19.93934

mape(test[,12],predict_dt)

## [1] 26.98542

mape(test[,12], predict_rf)

## [1] 24.413
```

3.1.2 Mean Squared Error (MSE)

Mean Squared Error(MSE) and Root Mean Squared Error (RMSE) can also be used as our metrics to evaluate the performance of our models. We can use `regr.eval()` function to get our RMSE, MSE, MAPE and MAE results. Based on these metrics we can evaluate which model gives good performance and better accuracy.

Based on our below evaluation, we see that `predict_lm_AIC` model gives better MSE results.

```
regr.eval(test[,12],predict_lm, stats = c("rmse", "mse", "mape", "mae"))
##          rmse          mse          mape          mae
## 8.100396e+02 6.561641e+05 2.016607e-01 6.103268e+02

regr.eval(test[,12],predict_lm_AIC, stats = c("rmse", "mse", "mape", "mae"))
##          rmse          mse          mape          mae
## 7.823221e+02 6.120279e+05 1.993934e-01 5.920271e+02

regr.eval(test[,12],predict_dt, stats = c("rmse", "mse", "mape", "mae"))
##          rmse          mse          mape          mae
## 1.005612e+03 1.011255e+06 2.698542e-01 7.234144e+02

regr.eval(test[,12],predict_rf, stats = c("rmse", "mse", "mape", "mae"))
##          rmse          mse          mape          mae
##   810.10792 656274.84090      0.24413    597.85694
```

3.2 Final Model Selection

From our model evaluation, we can see that stepwise fitting regression model on linear regression algorithm gives us low error rate and better results compared to decision trees and random forest.

Appendix A - R Code

```
##### Feature Engineering #####
data_df$season = as.factor(data_df$season)
levels(data_df$season) = c("Spring", "Summer", "Fall", "Winter")
data_df$yr = as.factor(data_df$yr)
levels(data_df$yr) = c("2011", "2012")
data_df$mnth = as.factor(data_df$mnth)
levels(data_df$mnth) = c("Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec")
data_df$holiday = as.factor(data_df$holiday)
levels(data_df$holiday) = c("No", "Yes")
data_df$weekday = as.factor(data_df$weekday)
levels(data_df$weekday) = c("Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat")
data_df$workingday = as.factor(data_df$workingday)
levels(data_df$workingday) = c("No", "Yes")
data_df$weathersit = as.factor(data_df$weathersit)
levels(data_df$weathersit) = c("Clear + Few Clouds", "Mist + Cloudy", "Light Snow + Rain + Thunderstorm")
data_df$dteday = as.POSIXct(data_df$dteday)
data_df$day = strftime(data_df$dteday, '%e')
data_df$day = as.factor(data_df$day)

##### Exploratory Data Analysis #####
#Analysis on dependent variable - Count
ggplot(data = data_df, aes(cnt))+
  geom_histogram()
```

```
#Analysis of Independent variables vs dependent variable
#Season vs Rental counts
g1 = ggplot(data = data_df, aes(data_df$season, data_df$cnt, fill = season))
  +
  geom_boxplot(show.legend = FALSE)+
  xlab("Season")+
  ylab("Rental Count")+
  labs(title = "Bike Rental counts during seasons")+
  theme(plot.title = element_text(hjust=0.5))

g2 = ggplot(data = data_df, aes(data_df$yr, data_df$cnt, fill = data_df$yr))
  +
  geom_boxplot(show.legend = FALSE)+
  xlab("Year")+
  ylab("Rental Count")+
  labs(title = "Year wise Bike Rental counts")+
  theme(plot.title = element_text(hjust=0.5))
```

```

g3 = ggplot(data = data_df, aes(data_df$mnth, data_df$cnt, fill = data_df$mnth))+
  geom_boxplot(show.legend = FALSE)+
  xlab("Month")+
  ylab("Rental Count")+
  labs(title = "Month wise Bike Rental counts")+
  theme(plot.title = element_text(hjust=0.5))

g4 = ggplot(data = data_df, aes(data_df$day, data_df$cnt, fill = data_df$day))+
  geom_boxplot(show.legend = FALSE)+
  xlab("Day")+
  ylab("Rental Count")+
  labs(title = "Day wise Bike Rental counts")+
  theme(plot.title = element_text(hjust=0.5))

g5 = ggplot(data = data_df, aes(data_df$holiday, data_df$cnt, fill = data_df$holiday))+
  geom_boxplot(show.legend = FALSE)+
  xlab("Holiday")+
  ylab("Rental Count")+
  labs(title = "Bike Rental counts during Holidays")+
  theme(plot.title = element_text(hjust=0.5))

g6 = ggplot(data = data_df, aes(data_df$weekday, data_df$cnt, fill = data_df$weekday))+
  geom_boxplot(show.legend = FALSE)+
  xlab("Weekday")+
  ylab("Rental Count")+
  labs(title = "Bike Rental counts during weekdays")+
  theme(plot.title = element_text(hjust=0.5))

g7 = ggplot(data = data_df, aes(data_df$workingday, data_df$cnt, fill = data_df$workingday))+
  geom_boxplot(show.legend = FALSE)+
  xlab("Workingday")+
  ylab("Count")+
  labs(title = "Bike Rental counts on a workingday")+
  theme(plot.title = element_text(hjust=0.5))

g8 = ggplot(data = data_df, aes(data_df$weathersit, data_df$cnt, fill = data_df$weathersit))+
  geom_boxplot(show.legend = FALSE)+
  xlab("Weather Condition")+
  ylab("Count")+
  labs(title = "Bike Rental counts based on weather conditions")+
  theme(plot.title = element_text(hjust=0.5))

```



```
gridExtra::grid.arrange(g1,g2,g3,g4,g5,g6,g7,g8,ncol=2)
```

```
#Distribution of Continuous variables
```

```
Hg1 = ggplot(data_df, aes(x = casual))+  
  geom_histogram()+  
  xlab("Casual Users")
```

```
Hg2 = ggplot(data_df, aes(x = windspeed))+  
  geom_histogram()+  
  xlab("WindSpeed")
```

```
Hg3 = ggplot(data_df, aes(x = temp))+  
  geom_histogram()+  
  xlab("Normalized Temperature")
```

```
Hg4 = ggplot(data_df, aes(x = atemp))+  
  geom_histogram()+  
  xlab("Normalized feeling temperature")
```

```
Hg5 = ggplot(data_df, aes(x = hum))+  
  geom_histogram()+  
  xlab("Humidity")
```

```
Hg6 = ggplot(data_df, aes(x = registered))+  
  geom_histogram()+  
  xlab("Registered Users")
```

```
gridExtra::grid.arrange(Hg1, Hg2,Hg3,Hg4,Hg5,Hg6)
```

```
##### Missing Values Analysis #####
```

```
Missing_val = data.frame(apply(data_df, 2, function(x){sum(is.na(x))}))
```

```
Missing_val$Columns = row.names(Missing_val)
```

```
row.names(Missing_val) = NULL
```

```
names(Missing_val)[1] = "Count"
```

```
#Display Missing Value counts - There are no missing values
```

```
Missing_val = Missing_val[c(2,1)]
```

```
##### Outlier Analysis #####
```

```
#Check for numeric variables and store them in numeric_index
```

```
numeric_index = sapply(data_df, is.numeric)
```

```
numeric_data = data_df[,numeric_index]
```

```
boxplot.stats(numeric_data$temp)$out
```

```
## numeric(0)
```

```
boxplot.stats(numeric_data$atemp)$out
```

```

## numeric(0)

  boxplot.stats(numeric_data$hum)$out
## [1] 0.187917 0.000000

  boxplot.stats(numeric_data$windspeed)$out
## [1] 0.417908 0.507463 0.385571 0.388067 0.422275 0.415429 0.409212
## [8] 0.421642 0.441563 0.414800 0.386821 0.398008 0.407346

  boxplot.stats(numeric_data$casual)$out
## [1] 2355 2282 3065 2418 2521 2397 3155 2469 2301 2347 3252 2795 2846 2541
## [15] 2496 2622 3410 2704 2855 3283 2557 2795 2494 2708 2963 2634 2657 2551
## [29] 2562 2355 2544 2345 2827 2352 2613 2570 3160 2512 2454 2589 3031 2806
## [43] 2643 2290

  boxplot.stats(numeric_data$registered)$out
## integer(0)

  numeric_data = data_df[,c("windspeed", "casual", "hum")]
  cnames = colnames(numeric_data)
  for(i in 1:length(cnames)){
    assign(paste0("Gn",i),
      ggplot(data_df, aes_string(y = cnames[i]))+
        stat_boxplot(geom = "errorbar", width = 0.5) +
        geom_boxplot(outlier.color = "red", outlier.shape = 18, outlier.
size = 2))
  }

  gridExtra::grid.arrange(Gn1,Gn2,Gn3,ncol=3)

```

```

##### Feature Selection #####
#Remove Casual and Registered from the dataset as this sums up total counts
#Remove instant variable and dteday as they do not give much information
data_df = data_df[, -c(1,2,14,15)]

#Correlation Analysis
numeric_index = sapply(data_df, is.numeric)
numeric_data = data_df[,numeric_index]

train_cor = cor(numeric_data)
corrplot(train_cor, method = 'color', addCoef.col = 'black')

```

```

#Feature selection through Random Forest
RF_model = randomForest(cnt ~., data = data_df, ntree = 100, importance = T

```

```

RUE)
pd = as.data.frame(importance(RF_model, type = 1))
pd$columns = row.names(pd)
row.names(pd) = NULL
pd = pd[order(-pd$`%IncMSE`),]
pd = pd[,c(2,1)]
View(pd)

#Plot features with imporatance from High to Low
ggplot(pd, aes(x = reorder(columns, -pd$`%IncMSE`), y = pd$`%IncMSE`, fill
= 'blue'))+
  geom_bar(stat = 'identity', show.legend = FALSE)+
  xlab("Variable Names")+
  ylab("Feature Importance Rate (%IncMSE)")+
  ggtitle("Features with Importance value")+
  theme(plot.title = element_text(hjust=0.5))

##### ModelDevelopment#####
#Split data
set.seed(1234)
train_index = createDataPartition(data_df$cnt, p = 0.8, list = F)
train = data_df[train_index,]
test = data_df[-train_index,]

#Check for multicollinearity
#library(usdm)
vif(numeric_data[, -5])

vifcor(numeric_data[, -5], th = 0.9)

#Multicollinearity issues are found between temp and atemp features

##### Linear Regression Modelling #####
lm_model = lm(data = train, cnt~ .)
summary(lm_model)

#Finding error metrics
#Calculating MAPE - y= Actual value, yhat = Predicted values
mape = function(y,yhat){
  mean(abs((y-yhat)/y))*100
}

mape(test[,12], predict_lm)

#Alternate method - regr.eval() from DMwR Library gives completed evaluatio
n of regression model
regr.eval(test[,12],predict_lm, stats = c("rmse", "mse", "mape", "mae"))

```

```

#Since we have multicollinearity issue between temp and atemp features we will regularize it
#Feature modeling using Stepwise modeling selection - both forward and backward
lm_AIC = stepAIC(lm_model, direction = 'both')

predict_lm_AIC = predict(lm_AIC, test[, -12])
summary(predict_lm_AIC)

mape(test[, 12], predict_lm_AIC)

regr.eval(test[, 12], predict_lm_AIC, stats = c("rmse", "mse", "mape", "mae"))

##### Decision Tree #####
dt_model = rpart(cnt ~., data = train, method = 'anova')

predict_dt = predict(dt_model, test[, -12])
summary(predict_dt)

mape(test[, 12], predict_dt)

regr.eval(test[, 12], predict_dt, stats = c("rmse", "mse", "mape", "mae"))

##### Random Forest #####
set.seed(3210)
rf_model = randomForest(cnt ~., data = train, ntree = 100)

predict_rf = predict(rf_model, test[, -12])
summary(predict_rf)

mape(test[, 12], predict_rf)
regr.eval(test[, 12], predict_rf, stats = c("rmse", "mse", "mape", "mae"))

```